

# PALLAV NAG

## 20BCE2023

### Assignment 2

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [2]: df=pd.read_csv('titanic.csv')
```

```
In [3]: df.head()
```

```
Out[3]:
```

	survived	pclass	sex	age	sibsp	parch	fare	embarked	class	who	adult_male	deck
0	0	3	male	22.0	1	0	7.2500	S	Third	man	True	Na
1	1	1	female	38.0	1	0	71.2833	C	First	woman	False	
2	1	3	female	26.0	0	0	7.9250	S	Third	woman	False	Na
3	1	1	female	35.0	1	0	53.1000	S	First	woman	False	
4	0	3	male	35.0	0	0	8.0500	S	Third	man	True	Na

```
In [4]: df.tail()
```

```
Out[4]:
```

	survived	pclass	sex	age	sibsp	parch	fare	embarked	class	who	adult_male	deck
886	0	2	male	27.0	0	0	13.00	S	Second	man	True	
887	1	1	female	19.0	0	0	30.00	S	First	woman	False	
888	0	3	female	NaN	1	2	23.45	S	Third	woman	False	
889	1	1	male	26.0	0	0	30.00	C	First	man	True	
890	0	3	male	32.0	0	0	7.75	Q	Third	man	True	

```
In [6]: df.columns
```

```
Out[6]: Index(['survived', 'pclass', 'sex', 'age', 'sibsp', 'parch', 'fare',
               'embarked', 'class', 'who', 'adult_male', 'deck', 'embark_town',
               'alive', 'alone'],
              dtype='object')
```

```
In [7]: df.shape
```

```
Out[7]: (891, 15)
```

```
In [8]: df.describe()
```


```
Out[8]:
```

	survived	pclass	age	sibsp	parch	fare
count	891.000000	891.000000	714.000000	891.000000	891.000000	891.000000
mean	0.383838	2.308642	29.699118	0.523008	0.381594	32.204208
std	0.486592	0.836071	14.526497	1.102743	0.806057	49.693429
min	0.000000	1.000000	0.420000	0.000000	0.000000	0.000000
25%	0.000000	2.000000	20.125000	0.000000	0.000000	7.910400
50%	0.000000	3.000000	28.000000	0.000000	0.000000	14.454200
75%	1.000000	3.000000	38.000000	1.000000	0.000000	31.000000
max	1.000000	3.000000	80.000000	8.000000	6.000000	512.329200

## Univariate analysis

```
In [9]: plt.hist(df['age'])
```

```
Out[9]: (array([ 54.,  46., 177., 169., 118.,  70.,  45.,  24.,   9.,   2.]),  
array([ 0.42 ,  8.378, 16.336, 24.294, 32.252, 40.21 , 48.168, 56.126,  
        64.084, 72.042, 80.   ]),  
<BarContainer object of 10 artists>)
```




```
In [20]: sns.countplot(df.sex)
```

D:\anaconda\lib\site-packages\seaborn\\_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

```
warnings.warn(
```

```
Out[20]: <AxesSubplot:xlabel='sex', ylabel='count'>
```




```
In [40]: sns.countplot(df.who)
```

D:\anaconda\lib\site-packages\seaborn\\_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

```
warnings.warn(
```

```
Out[40]: <AxesSubplot:xlabel='who', ylabel='count'>
```




## Bivariate analysis

### Fares paid according to class booked

```
In [18]: plt.scatter(df['pclass'],df['fare'])
```

```
Out[18]: <matplotlib.collections.PathCollection at 0x17cd8fad430>
```



```
In [27]: df.groupby(by='survived').agg('mean')[['pclass', 'age']]
```

```
Out[27]:
```


	pclass	age
survived		
0	2.531876	30.626179
1	1.950292	28.343690

```
In [29]: sns.barplot('class', 'fare', data=df)
```

D:\anaconda\lib\site-packages\seaborn\\_decorators.py:36: FutureWarning: Pass the following variables as keyword args: x, y. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.


warnings.warn(

```
Out[29]: <AxesSubplot:xlabel='class', ylabel='fare'>
```




```
In [41]: sns.countplot(data=df, x='alive', hue='sex')
```

```
Out[41]: <AxesSubplot:xlabel='alive', ylabel='count'>
```



```
In [42]: plt.figure(figsize=(12,8))
sns.kdeplot(data=df, x= 'pclass',hue='alive',fill=True)
```


```
Out[42]: <AxesSubplot:xlabel='pclass', ylabel='Density'>
```



## multivariate

```
In [30]: hm=df.corr()
sns.heatmap(hm)
```

```
Out[30]: <AxesSubplot:>
```



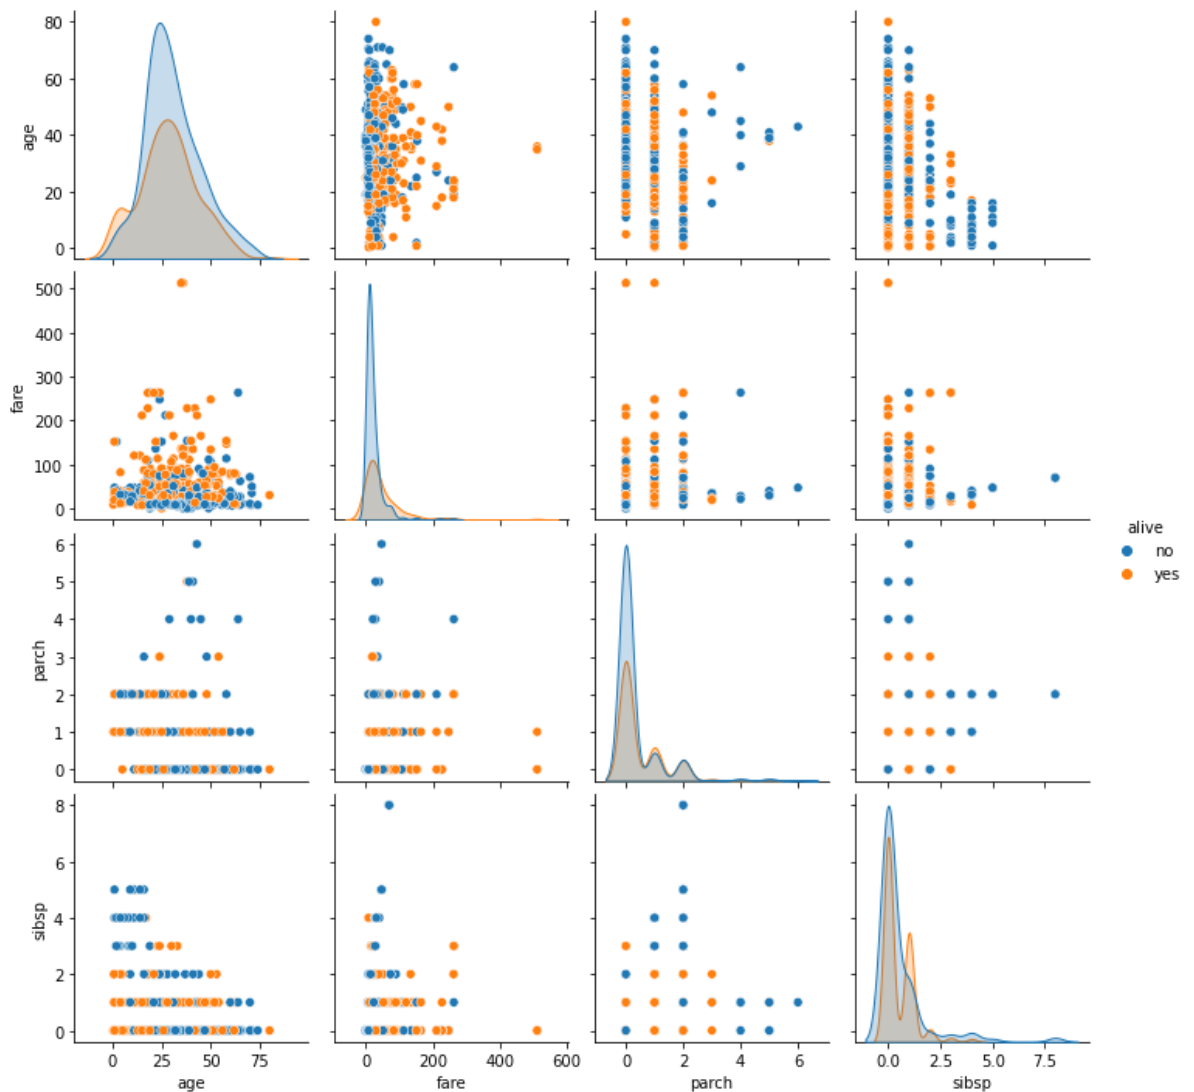
In [31]: hm

Out[31]:

	survived	pclass	age	sibsp	parch	fare	adult_male	alone
survived	1.000000	-0.338481	-0.077221	-0.035322	0.081629	0.257307	-0.557080	-0.203367
pclass	-0.338481	1.000000	-0.369226	0.083081	0.018443	-0.549500	0.094035	0.135207
age	-0.077221	-0.369226	1.000000	-0.308247	-0.189119	0.096067	0.280328	0.198270
sibsp	-0.035322	0.083081	-0.308247	1.000000	0.414838	0.159651	-0.253586	-0.584471
parch	0.081629	0.018443	-0.189119	0.414838	1.000000	0.216225	-0.349943	-0.583398
fare	0.257307	-0.549500	0.096067	0.159651	0.216225	1.000000	-0.182024	-0.271832
adult_male	-0.557080	0.094035	0.280328	-0.253586	-0.349943	-0.182024	1.000000	0.404744
alone	-0.203367	0.135207	0.198270	-0.584471	-0.583398	-0.271832	0.404744	1.000000

In [43]: sns.pairplot(data=df[['age', 'fare', 'parch', 'sibsp', 'alive']], hue='alive')

Out[43]: <seaborn.axisgrid.PairGrid at 0x17ce2a26970>



## descriptive statistics

In [32]: `df.mean()`

C:\Users\HP\AppData\Local\Temp\ipykernel\_12960\3698961737.py:1: FutureWarning: Dropping of nuisance columns in DataFrame reductions (with 'numeric\_only=None') is deprecated; in a future version this will raise TypeError. Select only valid columns before calling the reduction.

`df.mean()`

Out[32]:

survived	0.383838
pclass	2.308642
age	29.699118
sibsp	0.523008
parch	0.381594
fare	32.204208
adult_male	0.602694
alone	0.602694
dtype:	float64

In [33]: `df.median()`

C:\Users\HP\AppData\Local\Temp\ipykernel\_12960\530051474.py:1: FutureWarning: Dropping of nuisance columns in DataFrame reductions (with 'numeric\_only=None') is deprecated; in a future version this will raise TypeError. Select only valid columns before calling the reduction.

`df.median()`

```
Out[33]: survived      0.0000
         pclass       3.0000
         age         28.0000
         sibsp       0.0000
         parch       0.0000
         fare       14.4542
         adult_male   1.0000
         alone       1.0000
         dtype: float64
```

```
In [34]: df.mode()
```

```
Out[34]:
```

	survived	pclass	sex	age	sibsp	parch	fare	embarked	class	who	adult_male	deck	eml
0	0	3	male	24.0	0	0	8.05	S	Third	man	True	C	So

```
In [35]: df.skew()
```

C:\Users\HP\AppData\Local\Temp\ipykernel\_12960\1665899112.py:1: FutureWarning: Dropping of nuisance columns in DataFrame reductions (with 'numeric\_only=None') is deprecated; in a future version this will raise TypeError. Select only valid columns before calling the reduction.

```
df.skew()
```

```
Out[35]: survived      0.478523
         pclass     -0.630548
         age        0.389108
         sibsp      3.695352
         parch      2.749117
         fare      4.787317
         adult_male -0.420431
         alone     -0.420431
         dtype: float64
```

```
In [36]: df.kurt()
```

C:\Users\HP\AppData\Local\Temp\ipykernel\_12960\1257127604.py:1: FutureWarning: Dropping of nuisance columns in DataFrame reductions (with 'numeric\_only=None') is deprecated; in a future version this will raise TypeError. Select only valid columns before calling the reduction.

```
df.kurt()
```

```
Out[36]: survived     -1.775005
         pclass     -1.280015
         age        0.178274
         sibsp     17.880420
         parch      9.778125
         fare     33.398141
         adult_male -1.827345
         alone     -1.827345
         dtype: float64
```

```
In [37]: df.max()
```

C:\Users\HP\AppData\Local\Temp\ipykernel\_12960\1151452817.py:1: FutureWarning: Dropping of nuisance columns in DataFrame reductions (with 'numeric\_only=None') is deprecated; in a future version this will raise TypeError. Select only valid columns before calling the reduction.

```
df.max()
```



```
Out[37]: survived      1
         pclass        3
         sex          male
         age         80.0
         sibsp        8
         parch        6
         fare      512.3292
         class       Third
         who         woman
         adult_male   True
         alive        yes
         alone        True
         dtype: object
```

```
In [38]: df.min()
```

C:\Users\HP\AppData\Local\Temp\ipykernel\_12960\3962516015.py:1: FutureWarning: Dropping of nuisance columns in DataFrame reductions (with 'numeric\_only=None') is deprecated; in a future version this will raise TypeError. Select only valid columns before calling the reduction.

```
df.min()
```

```
Out[38]: survived      0
         pclass        1
         sex         female
         age         0.42
         sibsp        0
         parch        0
         fare         0.0
         class       First
         who         child
         adult_male   False
         alive        no
         alone        False
         dtype: object
```

## Handling missing data

```
In [45]: df.isnull().sum()
```

```
Out[45]: survived      0
         pclass        0
         sex          0
         age         177
         sibsp        0
         parch        0
         fare         0
         embarked     2
         class        0
         who          0
         adult_male   0
         deck        688
         embark_town  2
         alive        0
         alone        0
         dtype: int64
```

```
In [46]: df1 = pd.read_csv('titanic.csv')
         df2= pd.read_csv('titanic.csv')
         df3 = pd.read_csv('titanic.csv')
         df4= pd.read_csv('titanic.csv')
         df5=pd.read_csv('titanic.csv')
```

```
In [47]: #METHOD:1 DROP ALL ROWS WITH MISSING VALUE
df1.dropna(inplace= True)
df1.isnull().sum()
```

```
Out[47]: survived      0
pclass      0
sex         0
age         0
sibsp      0
parch      0
fare        0
embarked    0
class       0
who         0
adult_male  0
deck        0
embark_town 0
alive       0
alone       0
dtype: int64
```

```
In [48]: df1
```

```
Out[48]:
```

	survived	pclass	sex	age	sibsp	parch	fare	embarked	class	who	adult_male
1	1	1	female	38.0	1	0	71.2833	C	First	woman	False
3	1	1	female	35.0	1	0	53.1000	S	First	woman	False
6	0	1	male	54.0	0	0	51.8625	S	First	man	True
10	1	3	female	4.0	1	1	16.7000	S	Third	child	False
11	1	1	female	58.0	0	0	26.5500	S	First	woman	False
...	...	...	...	...	...	...	...	...	...	...	...
871	1	1	female	47.0	1	1	52.5542	S	First	woman	False
872	0	1	male	33.0	0	0	5.0000	S	First	man	True
879	1	1	female	56.0	0	1	83.1583	C	First	woman	False
887	1	1	female	19.0	0	0	30.0000	S	First	woman	False
889	1	1	male	26.0	0	0	30.0000	C	First	man	True

182 rows × 15 columns

```
In [49]: df2['age'].replace(np.NaN,df2['age'].mean()).head(30)
```

```
Out[49]: 0      22.000000
          1      38.000000
          2      26.000000
          3      35.000000
          4      35.000000
          5      29.699118
          6      54.000000
          7       2.000000
          8      27.000000
          9      14.000000
         10       4.000000
         11      58.000000
         12      20.000000
         13      39.000000
         14      14.000000
         15      55.000000
         16       2.000000
         17      29.699118
         18      31.000000
         19      29.699118
         20      35.000000
         21      34.000000
         22      15.000000
         23      28.000000
         24       8.000000
         25      38.000000
         26      29.699118
         27      19.000000
         28      29.699118
         29      29.699118
          Name: age, dtype: float64
```

```
In [51]: df3['age'].replace(np.NaN,df3['age'].median()).head(30)
```

```
Out[51]: 0      22.0
          1      38.0
          2      26.0
          3      35.0
          4      35.0
          5      28.0
          6      54.0
          7       2.0
          8      27.0
          9      14.0
         10       4.0
         11      58.0
         12      20.0
         13      39.0
         14      14.0
         15      55.0
         16       2.0
         17      28.0
         18      31.0
         19      28.0
         20      35.0
         21      34.0
         22      15.0
         23      28.0
         24       8.0
         25      38.0
         26      28.0
         27      19.0
         28      28.0
         29      28.0
          Name: age, dtype: float64
```

```
In [54]: df4['age'].replace(np.NaN,24).head(30) #mode=24
```

```
Out[54]: 0    22.0
1    38.0
2    26.0
3    35.0
4    35.0
5    24.0
6    54.0
7     2.0
8    27.0
9    14.0
10    4.0
11    58.0
12    20.0
13    39.0
14    14.0
15    55.0
16     2.0
17    24.0
18    31.0
19    24.0
20    35.0
21    34.0
22    15.0
23    28.0
24     8.0
25    38.0
26    24.0
27    19.0
28    24.0
29    24.0
Name: age, dtype: float64
```

```
In [56]: df['deck']=df3['deck'].fillna('NO')
df['age']=df3['age'].replace(np.NaN,df3['age'].median())
```

```
In [57]: df
```

```
Out[57]:
```

	survived	pclass	sex	age	sibsp	parch	fare	embarked	class	who	adult_male
0	0	3	male	22.0	1	0	7.2500	S	Third	man	True
1	1	1	female	38.0	1	0	71.2833	C	First	woman	False
2	1	3	female	26.0	0	0	7.9250	S	Third	woman	False
3	1	1	female	35.0	1	0	53.1000	S	First	woman	False
4	0	3	male	35.0	0	0	8.0500	S	Third	man	True
...	...	...	...	...	...	...	...	...	...	...	...
886	0	2	male	27.0	0	0	13.0000	S	Second	man	True
887	1	1	female	19.0	0	0	30.0000	S	First	woman	False
888	0	3	female	28.0	1	2	23.4500	S	Third	woman	False
889	1	1	male	26.0	0	0	30.0000	C	First	man	True
890	0	3	male	32.0	0	0	7.7500	Q	Third	man	True

891 rows × 15 columns


# Handling Outliers

```
In [58]: sns.boxplot(df1['age'])
```

D:\anaconda\lib\site-packages\seaborn\\_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

```
warnings.warn(
```

```
Out[58]: <AxesSubplot:xlabel='age'>
```



```
In [59]: np.where(df1['age']>62)
```

```
Out[59]: (array([ 9, 15, 50, 85, 90, 128, 155], dtype=int64),)
```

```
In [60]: from scipy import stats
z = np.abs(stats.zscore(df['age']))
z
```

```
Out[60]: 0      0.565736
1      0.663861
2      0.258337
3      0.433312
4      0.433312
...
886    0.181487
887    0.796286
888    0.104637
889    0.258337
890    0.202762
Name: age, Length: 891, dtype: float64
```

```
In [61]: np.where(z > 2)
```

```
Out[61]: (array([ 7, 11, 16, 33, 43, 54, 78, 94, 96, 116, 119, 152, 164,
170, 172, 174, 183, 193, 195, 205, 232, 252, 261, 268, 275, 280,
297, 305, 326, 340, 348, 366, 374, 381, 386, 407, 438, 456, 467,
469, 479, 483, 487, 493, 530, 545, 555, 570, 587, 625, 626, 630,
642, 644, 647, 659, 672, 684, 694, 745, 755, 772, 788, 803, 824,
827, 829, 831, 851, 879], dtype=int64),)
```

```
In [62]: Q1 = np.percentile(df['age'], 25)
Q3 = np.percentile(df['age'], 75)
IQR = Q3 - Q1
```

```
In [63]: upper=Q3+1.5*IQR
upper_array=np.array(df['age']>=upper)
print("Upper Bound:",upper)

lower=Q1-1.5*IQR
lower_array=np.array(df['age']<=lower)
print("Lower Bound:",lower)
```

Upper Bound: 54.5  
Lower Bound: 2.5

## Checking for Categorical columns and perform encoding.

```
In [64]: categorical_columns = df.select_dtypes(include=['object']).columns
```

```
In [65]: for column in categorical_columns:
          print(column)
```

sex  
embarked  
class  
who  
deck  
embark\_town  
alive

```
In [66]: encoded_df = pd.get_dummies(df, columns=categorical_columns)
```

```
In [67]: print("\nEncoded dataframe:")
          print(encoded_df)
```

Encoded dataframe:

	survived	pclass	age	sibsp	parch	fare	adult_male	alone	\
0	0	3	22.0	1	0	7.2500	True	False	
1	1	1	38.0	1	0	71.2833	False	False	
2	1	3	26.0	0	0	7.9250	False	True	
3	1	1	35.0	1	0	53.1000	False	False	
4	0	3	35.0	0	0	8.0500	True	True	
..	...	...	...	...	...	...	...	...	
886	0	2	27.0	0	0	13.0000	True	True	
887	1	1	19.0	0	0	30.0000	False	True	
888	0	3	28.0	1	2	23.4500	False	False	
889	1	1	26.0	0	0	30.0000	True	True	
890	0	3	32.0	0	0	7.7500	True	True	

	sex_female	sex_male	...	deck_D	deck_E	deck_F	deck_G	deck_NO	\
0	0	1	...	0	0	0	0	1	
1	1	0	...	0	0	0	0	0	
2	1	0	...	0	0	0	0	1	
3	1	0	...	0	0	0	0	0	
4	0	1	...	0	0	0	0	1	
..	...	...	...	...	...	...	...	...	
886	0	1	...	0	0	0	0	1	
887	1	0	...	0	0	0	0	0	
888	1	0	...	0	0	0	0	1	
889	0	1	...	0	0	0	0	0	
890	0	1	...	0	0	0	0	1	

	embark_town_Chernbourg	embark_town_Queenstown	embark_town_Southampton	\
0	0	0	1	
1	1	0	0	
2	0	0	1	
3	0	0	1	
4	0	0	1	
..	...	...	...	
886	0	0	1	
887	0	0	1	
888	0	0	1	
889	1	0	0	
890	0	1	0	

	alive_no	alive_yes
0	1	0
1	0	1
2	0	1
3	0	1
4	1	0
..	...	...
886	1	0
887	0	1
888	1	0
889	0	1
890	1	0

[891 rows x 32 columns]

## Splitting the data into dependent and independent variables.

```
In [68]: X = df.drop('survived', axis=1) # Independent variables (all columns except 'surv'
y = df['survived'] # Dependent variable ('survived' column)
```



```
In [69]: print("Independent variables:")
print(X)
```

Independent variables:

	pclass	sex	age	sibsp	parch	fare	embarked	class	who	\
0	3	male	22.0	1	0	7.2500	S	Third	man	
1	1	female	38.0	1	0	71.2833	C	First	woman	
2	3	female	26.0	0	0	7.9250	S	Third	woman	
3	1	female	35.0	1	0	53.1000	S	First	woman	
4	3	male	35.0	0	0	8.0500	S	Third	man	
..	...	...	...	...	...	...	...	...	...	
886	2	male	27.0	0	0	13.0000	S	Second	man	
887	1	female	19.0	0	0	30.0000	S	First	woman	
888	3	female	28.0	1	2	23.4500	S	Third	woman	
889	1	male	26.0	0	0	30.0000	C	First	man	
890	3	male	32.0	0	0	7.7500	Q	Third	man	

	adult_male	deck	embark_town	alive	alone
0	True	NO	Southampton	no	False
1	False	C	Cherbourg	yes	False
2	False	NO	Southampton	yes	True
3	False	C	Southampton	yes	False
4	True	NO	Southampton	no	True
..	...	...	...	...	...
886	True	NO	Southampton	no	True
887	False	B	Southampton	yes	True
888	False	NO	Southampton	no	False
889	True	C	Cherbourg	yes	True
890	True	NO	Queenstown	no	True

[891 rows x 14 columns]

```
In [70]: print("\nDependent variable:")
print(y)
```

Dependent variable:

```
0      0
1      1
2      1
3      1
4      0
..
886    0
887    1
888    0
889    1
890    0
```

Name: survived, Length: 891, dtype: int64

## Scaling the independent variables

```
In [71]: dependent_variable = df['survived'] # Assuming 'survived' is the dependent variable
independent_variables = df.drop('survived', axis=1)
```

```
In [72]: scaled_independent_variables = (independent_variables - independent_variables.mean())
```

C:\Users\HP\AppData\Local\Temp\ipykernel\_12960\1149853797.py:1: FutureWarning: Dropping of nuisance columns in DataFrame reductions (with 'numeric\_only=None') is deprecated; in a future version this will raise TypeError. Select only valid columns before calling the reduction.

```
scaled_independent_variables = (independent_variables - independent_variables.mean()) / independent_variables.std()
```

C:\Users\HP\AppData\Local\Temp\ipykernel\_12960\1149853797.py:1: FutureWarning: Dropping of nuisance columns in DataFrame reductions (with 'numeric\_only=None') is deprecated; in a future version this will raise TypeError. Select only valid columns before calling the reduction.

```
scaled_independent_variables = (independent_variables - independent_variables.mean()) / independent_variables.std()
```

```
In [73]: print("Scaled independent variables:")
print(scaled_independent_variables)
```

Scaled independent variables:

	adult_male	age	alive	alone	class	deck	embark_town	embarked	\
0	0.811467	-0.565419	NaN	-1.230954	NaN	NaN	NaN	NaN	
1	-1.230954	0.663488	NaN	-1.230954	NaN	NaN	NaN	NaN	
2	-1.230954	-0.258192	NaN	0.811467	NaN	NaN	NaN	NaN	
3	-1.230954	0.433068	NaN	-1.230954	NaN	NaN	NaN	NaN	
4	0.811467	0.433068	NaN	0.811467	NaN	NaN	NaN	NaN	
..	...	...	...	...	...	...	...	...	
886	0.811467	-0.181385	NaN	0.811467	NaN	NaN	NaN	NaN	
887	-1.230954	-0.795839	NaN	0.811467	NaN	NaN	NaN	NaN	
888	-1.230954	-0.104579	NaN	-1.230954	NaN	NaN	NaN	NaN	
889	0.811467	-0.258192	NaN	0.811467	NaN	NaN	NaN	NaN	
890	0.811467	0.202648	NaN	0.811467	NaN	NaN	NaN	NaN	

	fare	parch	pclass	sex	sibsp	who
0	-0.502163	-0.473408	0.826913	NaN	0.432550	NaN
1	0.786404	-0.473408	-1.565228	NaN	0.432550	NaN
2	-0.488580	-0.473408	0.826913	NaN	-0.474279	NaN
3	0.420494	-0.473408	-1.565228	NaN	0.432550	NaN
4	-0.486064	-0.473408	0.826913	NaN	-0.474279	NaN
..	...	...	...	...	...	...
886	-0.386454	-0.473408	-0.369158	NaN	-0.474279	NaN
887	-0.044356	-0.473408	-1.565228	NaN	-0.474279	NaN
888	-0.176164	2.007806	0.826913	NaN	0.432550	NaN
889	-0.044356	-0.473408	-1.565228	NaN	-0.474279	NaN
890	-0.492101	-0.473408	0.826913	NaN	-0.474279	NaN

[891 rows x 14 columns]

## Splitting the data into training and testing

```
In [75]: from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
```

```
In [76]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
In [77]: # Print the shapes of the training and testing sets
print("Training set shape:", X_train.shape, y_train.shape)
print("Testing set shape:", X_test.shape, y_test.shape)
```

Training set shape: (712, 14) (712,)

Testing set shape: (179, 14) (179,)

```
In [ ]:
```